



# Intel Xeon<sup>®</sup> , Intel Xeon Phi<sup>™</sup> Coprocessor and KNL Architecture

Silvio Stanzani , Raphael Cóbe , Rogério Iope

UNESP - Núcleo de Computação Científica


[silvio@ncc.unesp.br](mailto:silvio@ncc.unesp.br) , [rmcobe@ncc.unesp.br](mailto:rmcobe@ncc.unesp.br) , [rogerio@ncc.unesp.br](mailto:rogerio@ncc.unesp.br)

# Module Outline

- **Intel Xeon and Intel® Xeon Phi™ Overview**
- Programming Model
- Vector Processing Unit
- Setting Expectations
- System Software and OS
- Application Environment


# Intel Xeon and Intel® Xeon Phi™ Overview

## Intel® Multicore Architecture




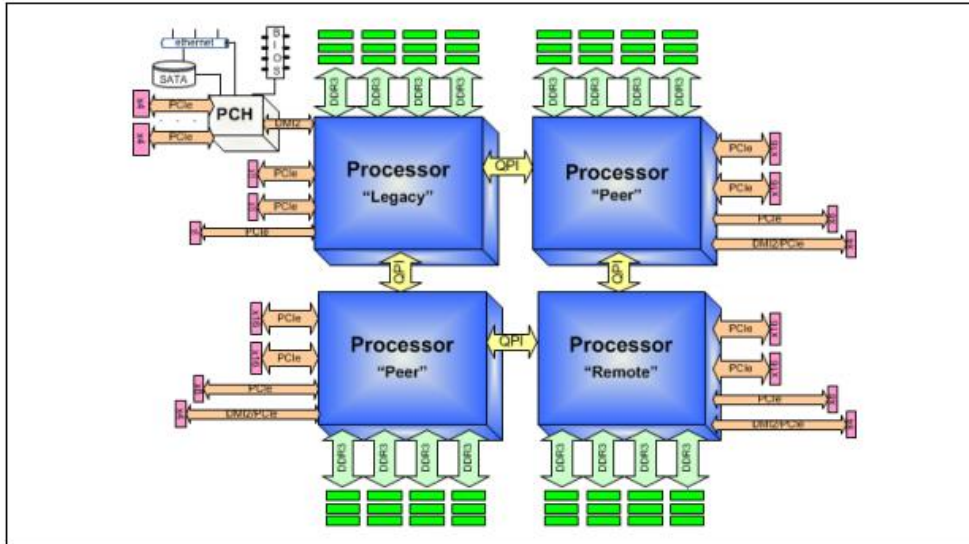
- ❖ Foundation of HPC Performance
- ❖ Suited for full scope of workloads
- ❖ Industry leading performance and performance/watt for serial & parallel workloads
- ❖ Focus on fast single core/thread performance with “moderate” number of cores

## Intel® Many Integrated Core Architecture



- ❖ Performance and performance/watt optimized for highly parallelized compute workloads
- ❖ Common software tools with Xeon enabling efficient application readiness and performance tuning
- ❖ IA extension to Manycore
- ❖ Many cores/threads with wide SIMD

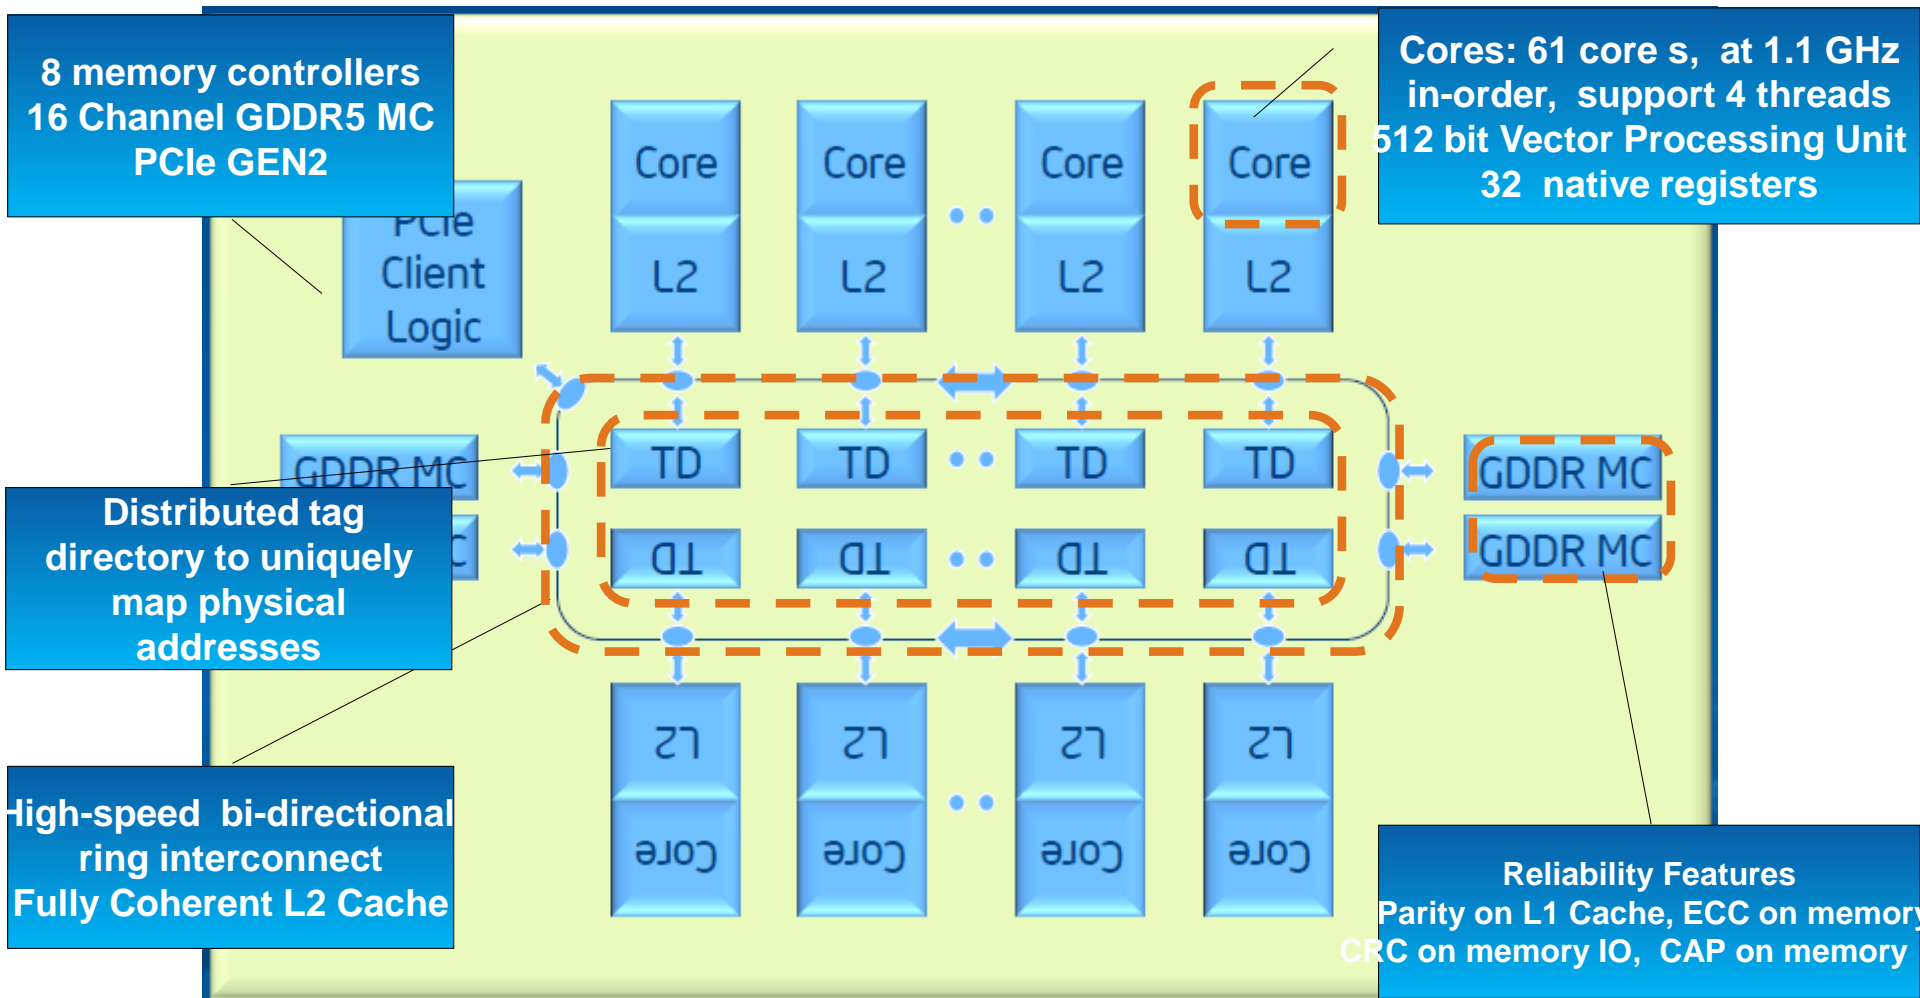
# Intel Xeon Architecture Overview



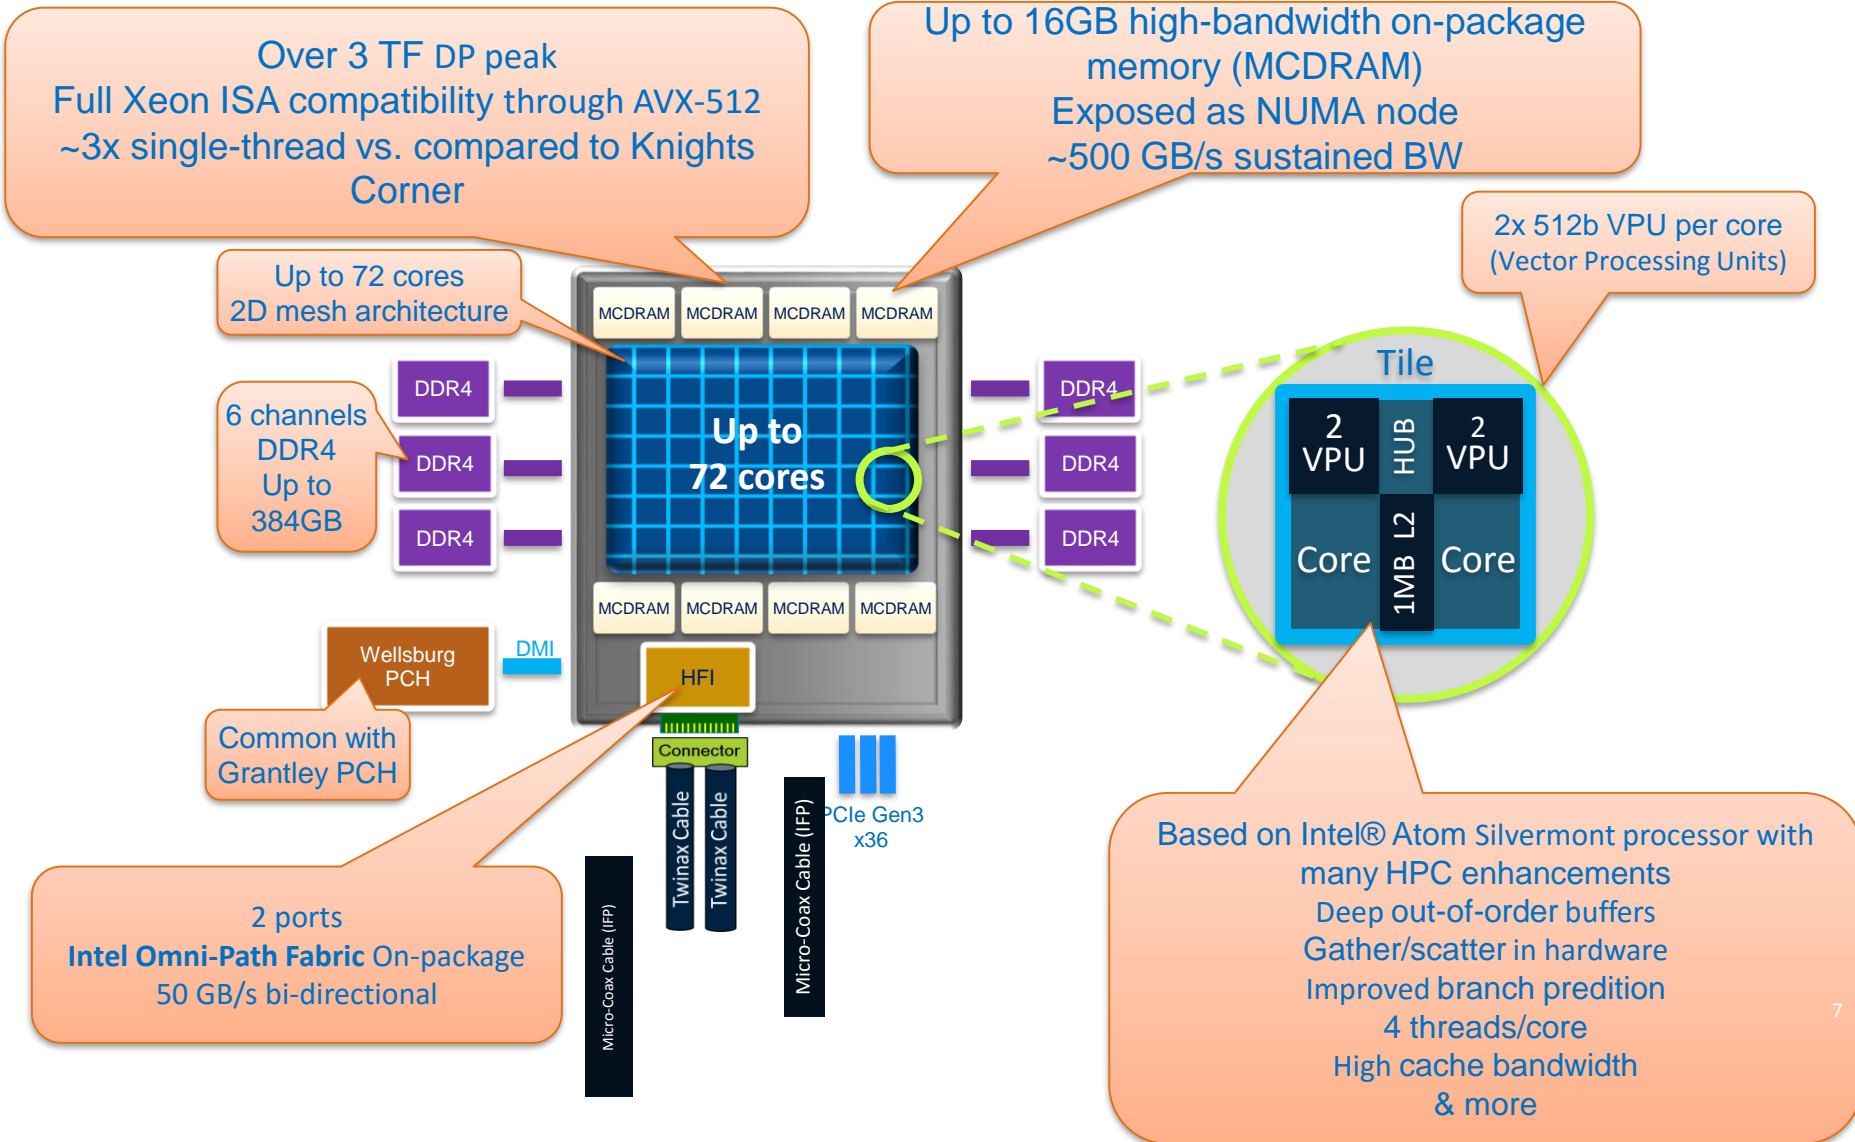
# Intel Xeon Architecture Overview

- Socket: mechanical component that provides mechanical and electrical connections between a microprocessor and a printed circuit board (PCB).
- QPI (Intel QuickPath Interconnect): high speed, packetized, point-to-point interconnection, that stitch together processors in distributed shared memory and integrated I/O platform architecture.

# Intel® Xeon Phi™ Architecture Overview



# Knights Landing (KNL)



# Knights Landing (KNL)

- KNL Knights Landing has 72 cores;
- Each one has an L1 cache;
- Pairs of cores are organized into tiles with a slice of the L2 cache symmetrically shared between the two cores;
- All caches are kept coherent;
- Two VPU (Vector Processing Units) per core;




# Cluster modes

## One single space address

### Hemisphere:

the tiles are divided into two parts called hemisphere


Node 0



### Quadrant:

tiles are divided into two parts called hemisphere or into four parts called quadrants

Node 0




# Cluster modes

Cache data are isolated in each sub numa domain


## SNC-2:

the tiles are  
divided into two  
Numa Nodes




## SNC-4:

the tiles are  
divided into two  
Numa Nodes




# Knights Landing Integrated On-Package Memory

Multi-Channel DRAM (High-bandwidth memory)



# Integrated On-Package Memory Usage Models

Model configurable at boot time and software exposed through NUMA



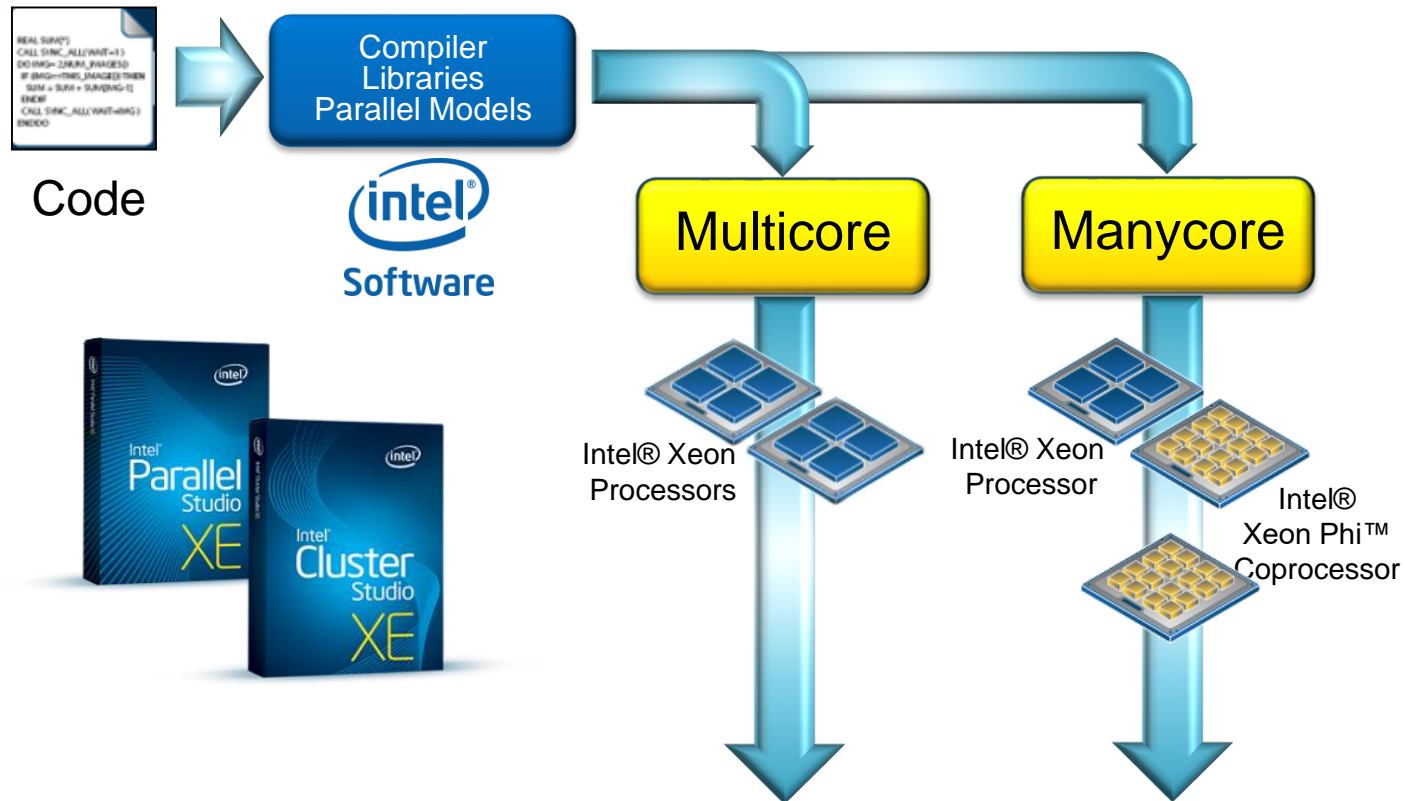
Split Options:  
25/75% or 50/50%

Cache Model	Flat Model	Hybrid Model
Hardware automatically manages the MCDRAM as a “L3 cache” between CPU and ext DDR memory	Manually manage how the app uses the integrated on-package memory and external DDR for peak perf	Harness the benefits of both Cache and Flat models by segmenting the integrated on-package memory
<ul style="list-style-type: none"><li>▪ App and/or data set is very large and will not fit into MCDRAM</li><li>▪ Unknown or unstructured memory access behavior</li></ul>	<ul style="list-style-type: none"><li>▪ App or portion of an app or data set that can be, or is needed to be “locked” into MCDRAM so it doesn’t get flushed out</li></ul>	<ul style="list-style-type: none"><li>▪ Need to “lock” in a relatively small portion of an app or data set via the Flat model</li><li>▪ Remaining MCDRAM can then be configured as Cache</li></ul>

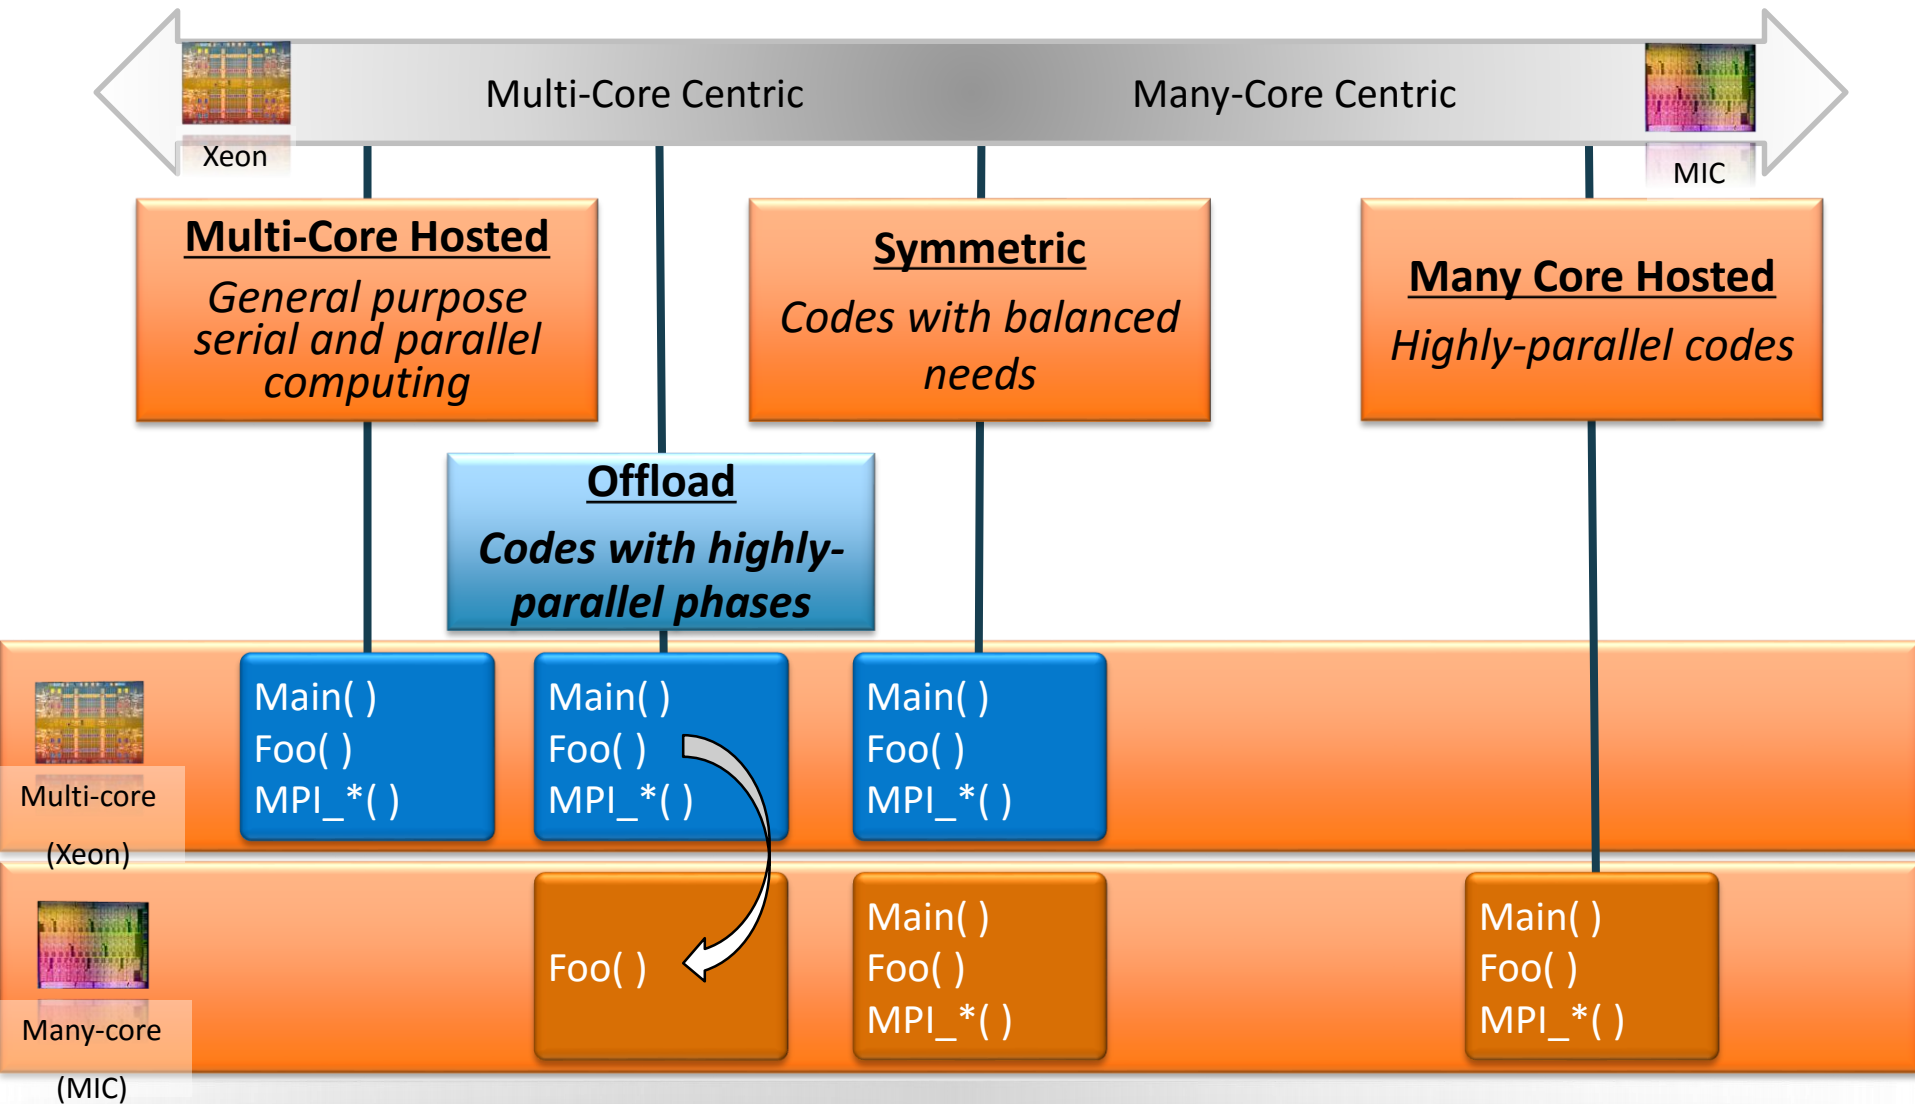
# Module Outline

- Intel Xeon and Intel® Xeon Phi™ Overview
- **Programming Model**
- Vector Processing Unit
- Setting Expectations
- System Software and OS
- Application Environment

# Programming Models



# Programming Models



# Module Outline

- Intel Xeon and Intel® Xeon Phi™ Overview
- Programming Model
- Vector Processing Unit
- **Setting Expectations**
- System Software and OS
- Application Environment



# Intel® MIC Programming Considerations

- Getting full performance from the Intel® MIC architecture requires both a high degree of parallelism and vectorization
  - Not all code can be written this way
  - Not all programs make sense on this architecture
- Intel® MIC is not an Intel® Xeon® processor
  - It specializes in running highly parallel and vectorized code.
  - New vector instruction set and 512-bit wide registers
  - Not optimized for processing serial code

# Intel® MIC Programming Considerations

- KNC comes with 8GB of memory
  - Only ~7GB is available to your program
    - ❑ The other ~1GB is used for data transfer and is accessible to your Intel® MIC Architecture code as buffers.
- Very short (low-latency) tasks not optimal for offload to the coprocessor
  - Costs that you need to amortize to make it worthwhile:
    - ❑ Cost of code and data transfer
    - ❑ Cost of process/thread creation
  - Fastest data transfers currently require careful data alignment
- This architecture is not optimized for serial performance

Thread setup  
and comm  
overhead

Off  
load

# Module Outline

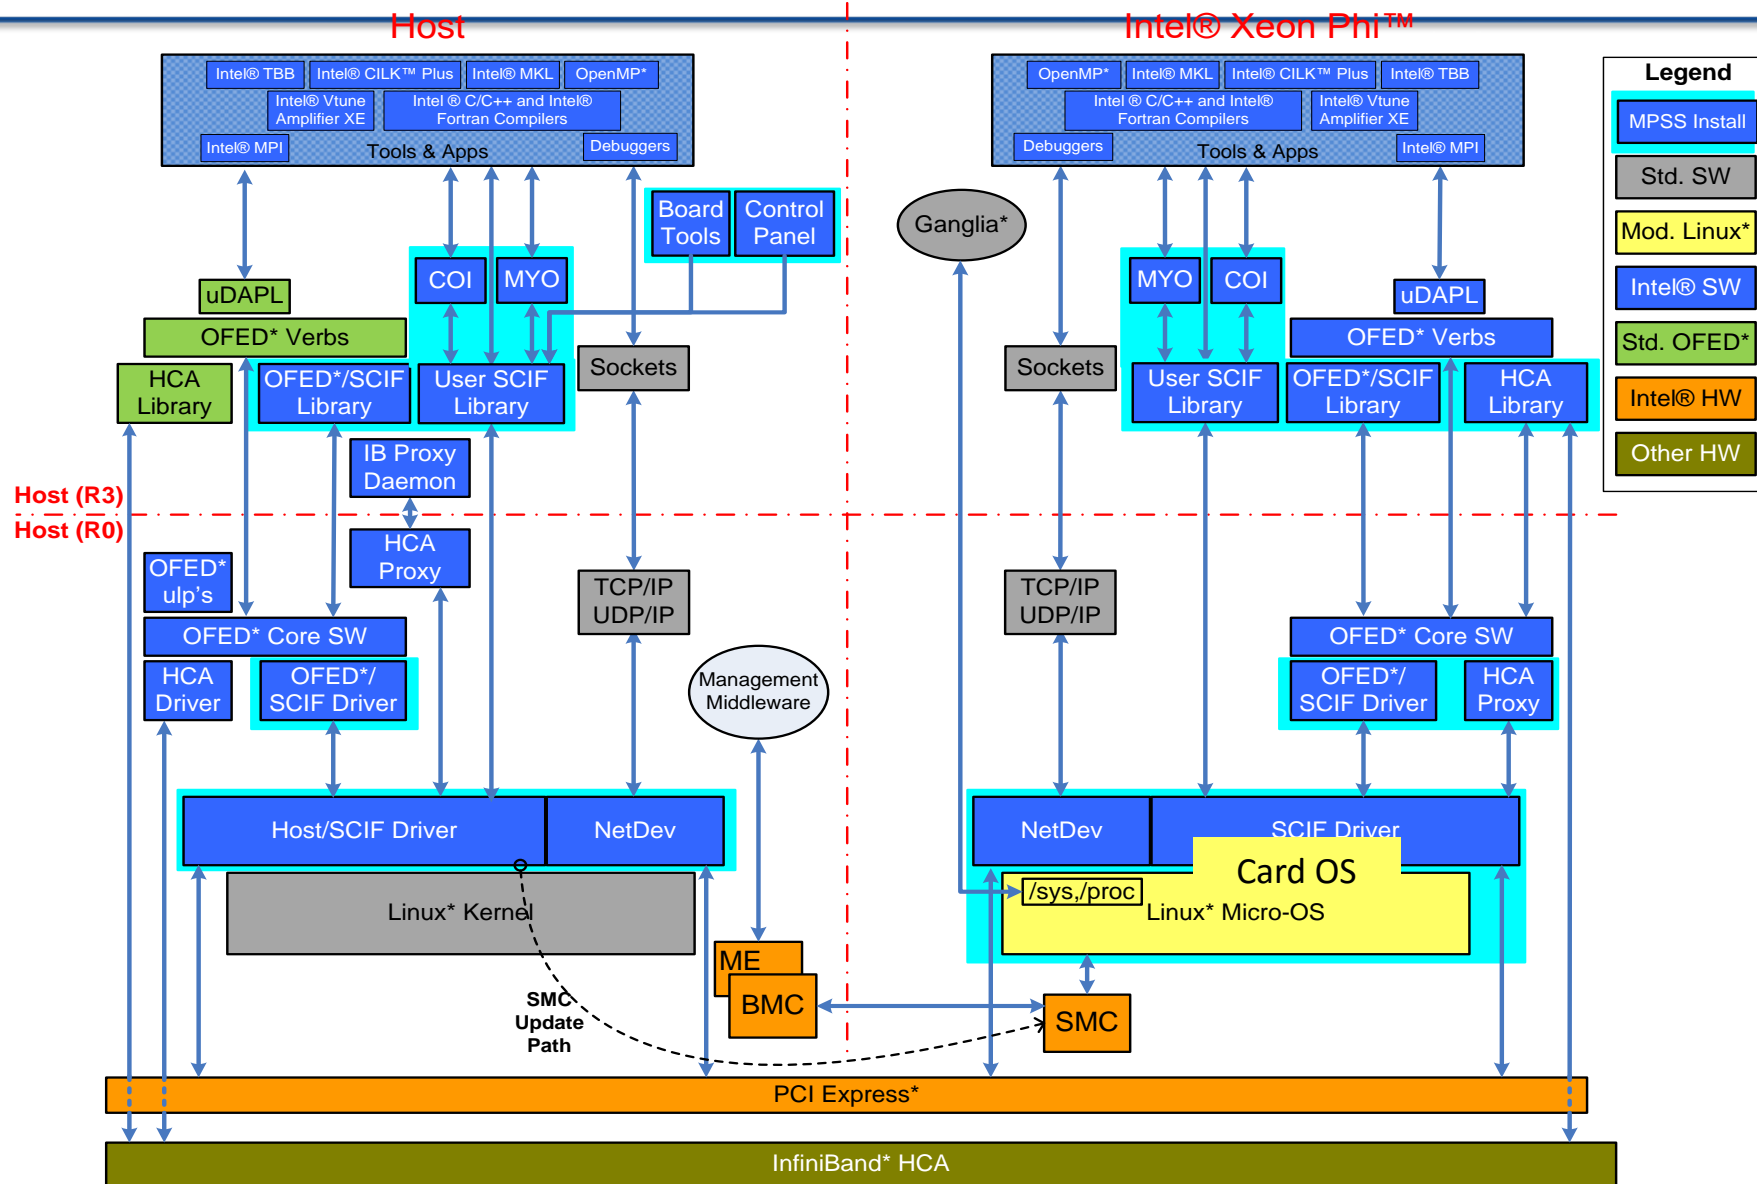
---

- Intel Xeon and Intel® Xeon Phi™ Overview
- Programming Model
- Vector Processing Unit
- Setting Expectations
- **System Software and OS**
- Application Environment

# Intel® Xeon Phi™ Coprocessor Arch – System SW Perspective

- Large SMP UMA machine – a set of x86 cores to manage
  - 4 threads and 32KB L1I/D, 512KB L2 per core
  - Supports loadable kernel modules
  - Standard Linux kernel from kernel.org
    - ❑ 2.6.38 in the most recent release
  - Completely Fair Scheduler (CFS), VM subsystem, File I/O
- Virtual Ethernet driver– supports NFS mounts from Intel® Xeon Phi™ Coprocessor
- New vector register state per thread for Intel® IMCI
  - Supports “Device Not Available” for Lazy save/restore
- Different ABI – uses vector registers for passing floats
  - Still uses the x86\_64 ABI for non-float parameter passing (rdi, rsi, rdx ..)

# System SW Environment



# Module Outline

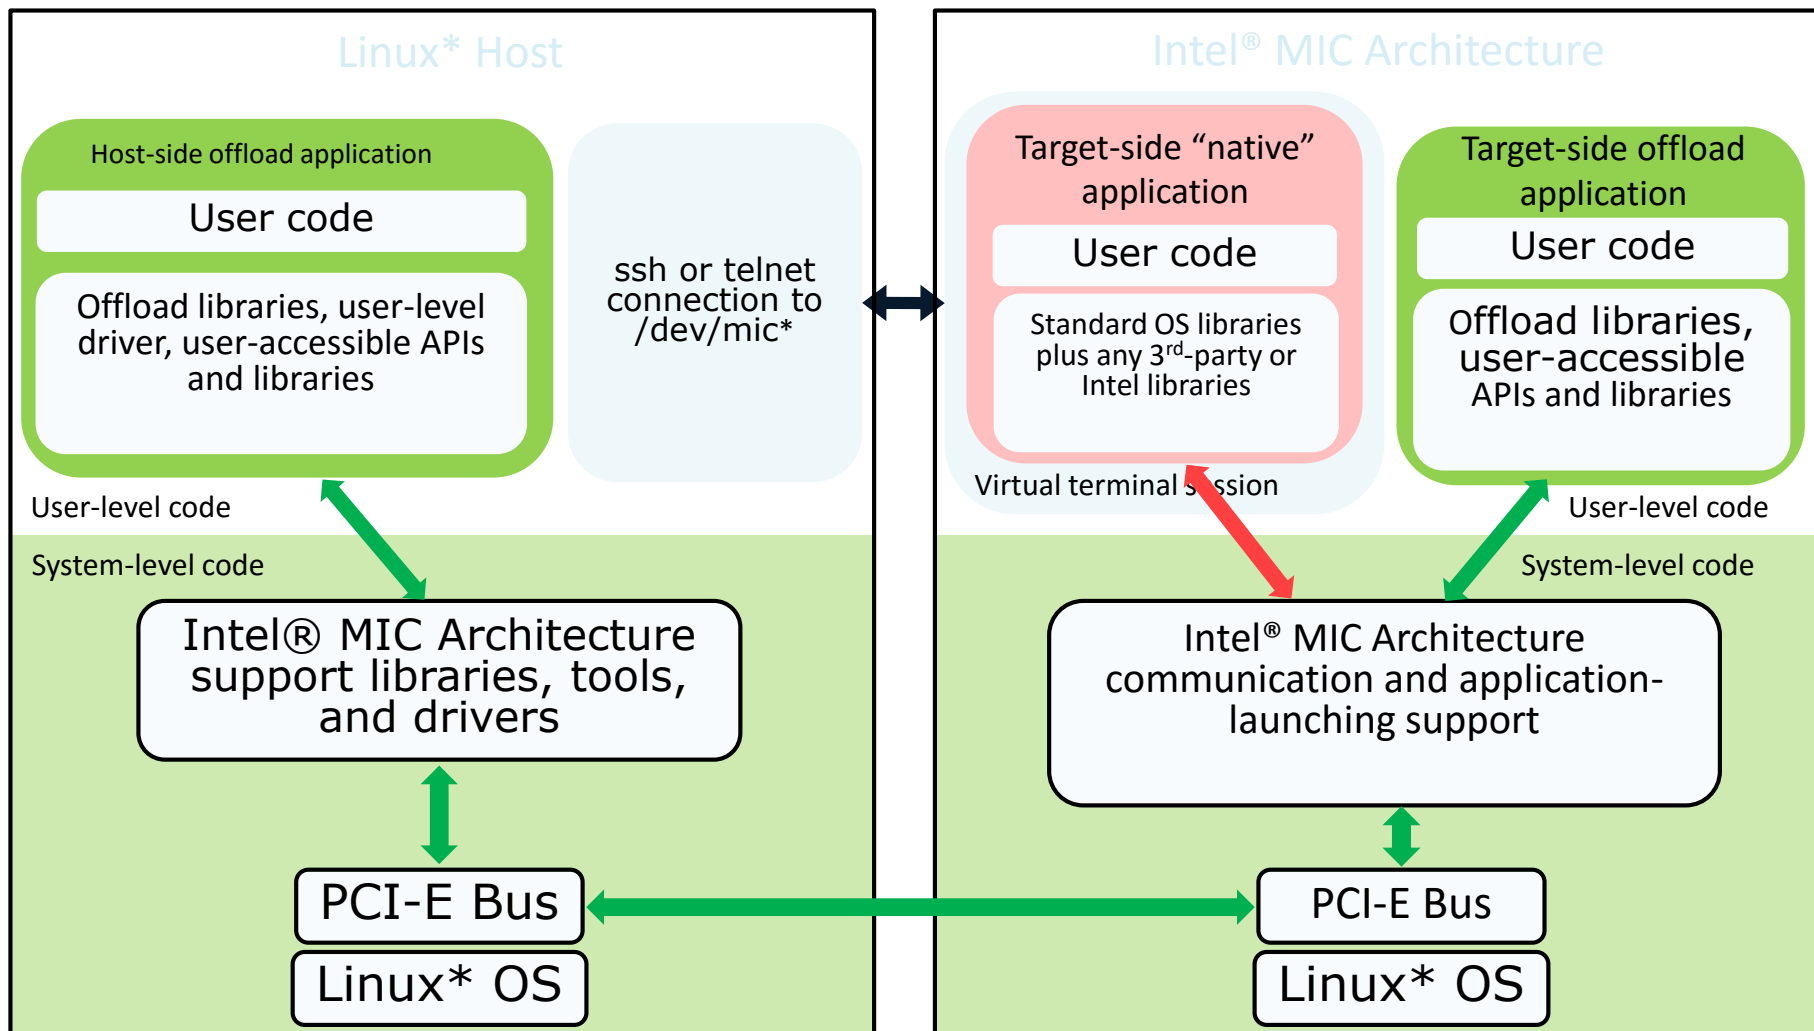
- Intel Xeon and Intel® Xeon Phi™ Coprocessor Overview
- Core and Vector Processing Unit
- Setting Expectations
- System Software and OS
- **Application Environment**

# Offload Programming Model Data Transfer

- Explicit distributed memory programming via MPI (or sockets)
- Two offload data transfer models are available:
  - Explicit Copy
    - ❑ Programmer designates variables that need to be copied between host and card in the offload directive
    - ❑ Syntax: Pragma/directive-based
    - ❑ C/C++ Example: `#pragma offload target(mic) in(data:length(size))`
    - ❑ Fortran Example: `!dir$ offload target(mic) in(a1:length(size))`
  - Implicit Copy
    - ❑ Programmer marks variables that need to be shared between host and card
    - ❑ The same variable can then be used in both host and coprocessor code
    - ❑ Runtime automatically maintains coherence at the beginning and end of offload statements
    - ❑ Syntax: keyword extensions based
    - ❑ Example: `_Cilk_shared double foo; _Offload func(y);`

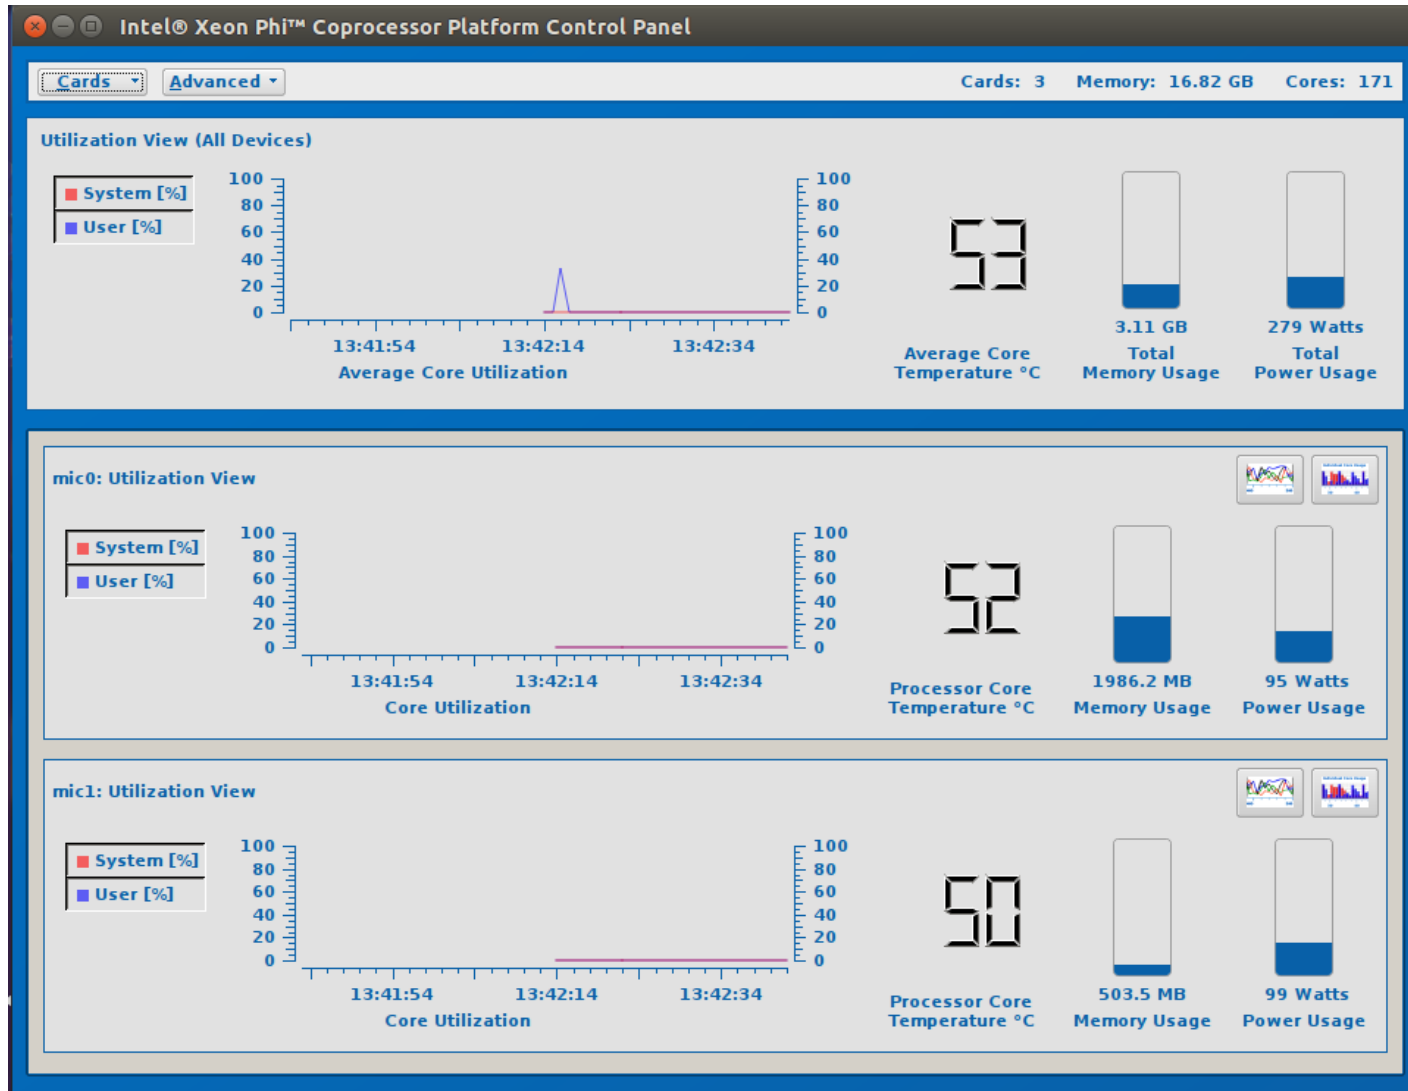
# Two Application Execution Environments

+Hybrid





# MICSMC



# Hello World

```
[silvio@phi01 ~]$ cat helloWorld.c
#include <stdio.h>

#include <unistd.h>

int main() {
    printf("Hello world! I have %ld logical cores.\n", sysconf(_SC_NPROCESSORS_ONLN ));
}
[silvio@phi01 ~]$ icc helloWorld.c -o helloWorld
[silvio@phi01 ~]$ ./helloWorld
Hello world! I have 32 logical cores.
[silvio@phi01 ~]$ icc helloWorld.c -o helloWorld.mic -mmic
[silvio@phi01 ~]$ ssh mic0
silvio@phi01-mic0:~$ ./helloWorld.mic
Hello world! I have 228 logical cores.
silvio@phi01-mic0:~$ █
```

