



# Intel Xeon<sup>®</sup>, Intel Xeon Phi<sup>™</sup> Coprocessor and KNL Architecture

Silvio Stanzani , Raphael Cóbe , Rogério Iope

UNESP - Núcleo de Computação Científica

[silvio@ncc.unesp.br](mailto:silvio@ncc.unesp.br) , [rmcobe@ncc.unesp.br](mailto:rmcobe@ncc.unesp.br) ,  
[rogerio@ncc.unesp.br](mailto:rogerio@ncc.unesp.br)

# Module Outline

---

- Intel Xeon and Intel® Xeon Phi™
- System Software and OS
- Intel KNL
- Programming Model


# Module Outline

---

- **Intel Xeon and Intel® Xeon Phi™**
- System Software and OS
- Intel KNL
- Programming Model


# Intel Xeon and Intel® Xeon Phi™ Overview

## Intel® Multicore Architecture



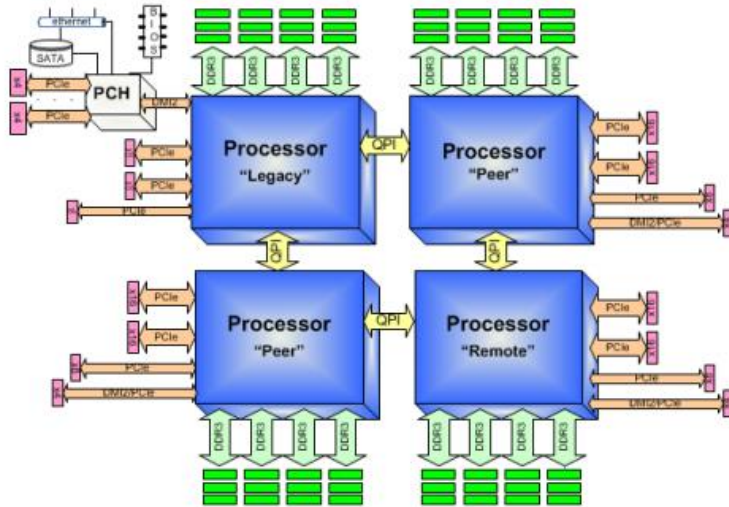
- ❖ Foundation of HPC Performance
- ❖ Suited for full scope of workloads
- ❖ Focus on fast single core/thread performance with “moderate” number of cores

## Intel® Many Integrated Core Architecture

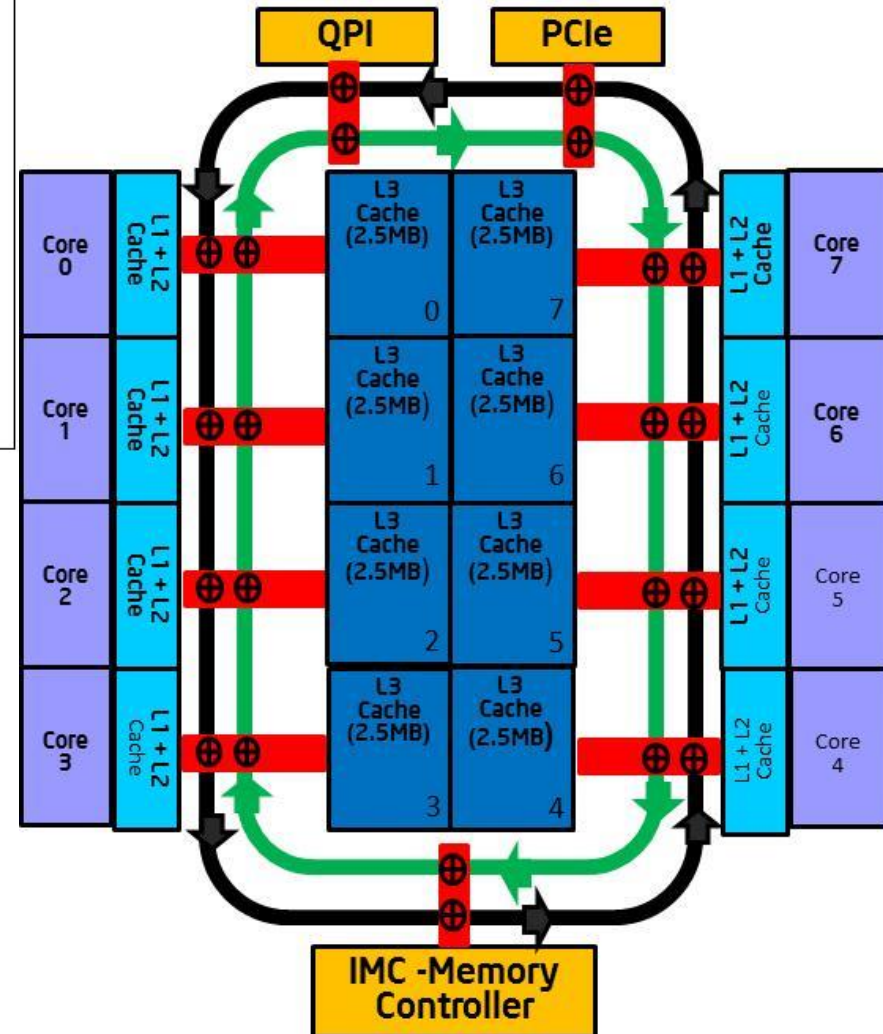


- ❖ Performance and performance/watt optimized for highly parallelized compute workloads
- ❖ IA extension to Manycore
- ❖ Many cores/threads with wide SIMD

# Intel Xeon Architecture Overview

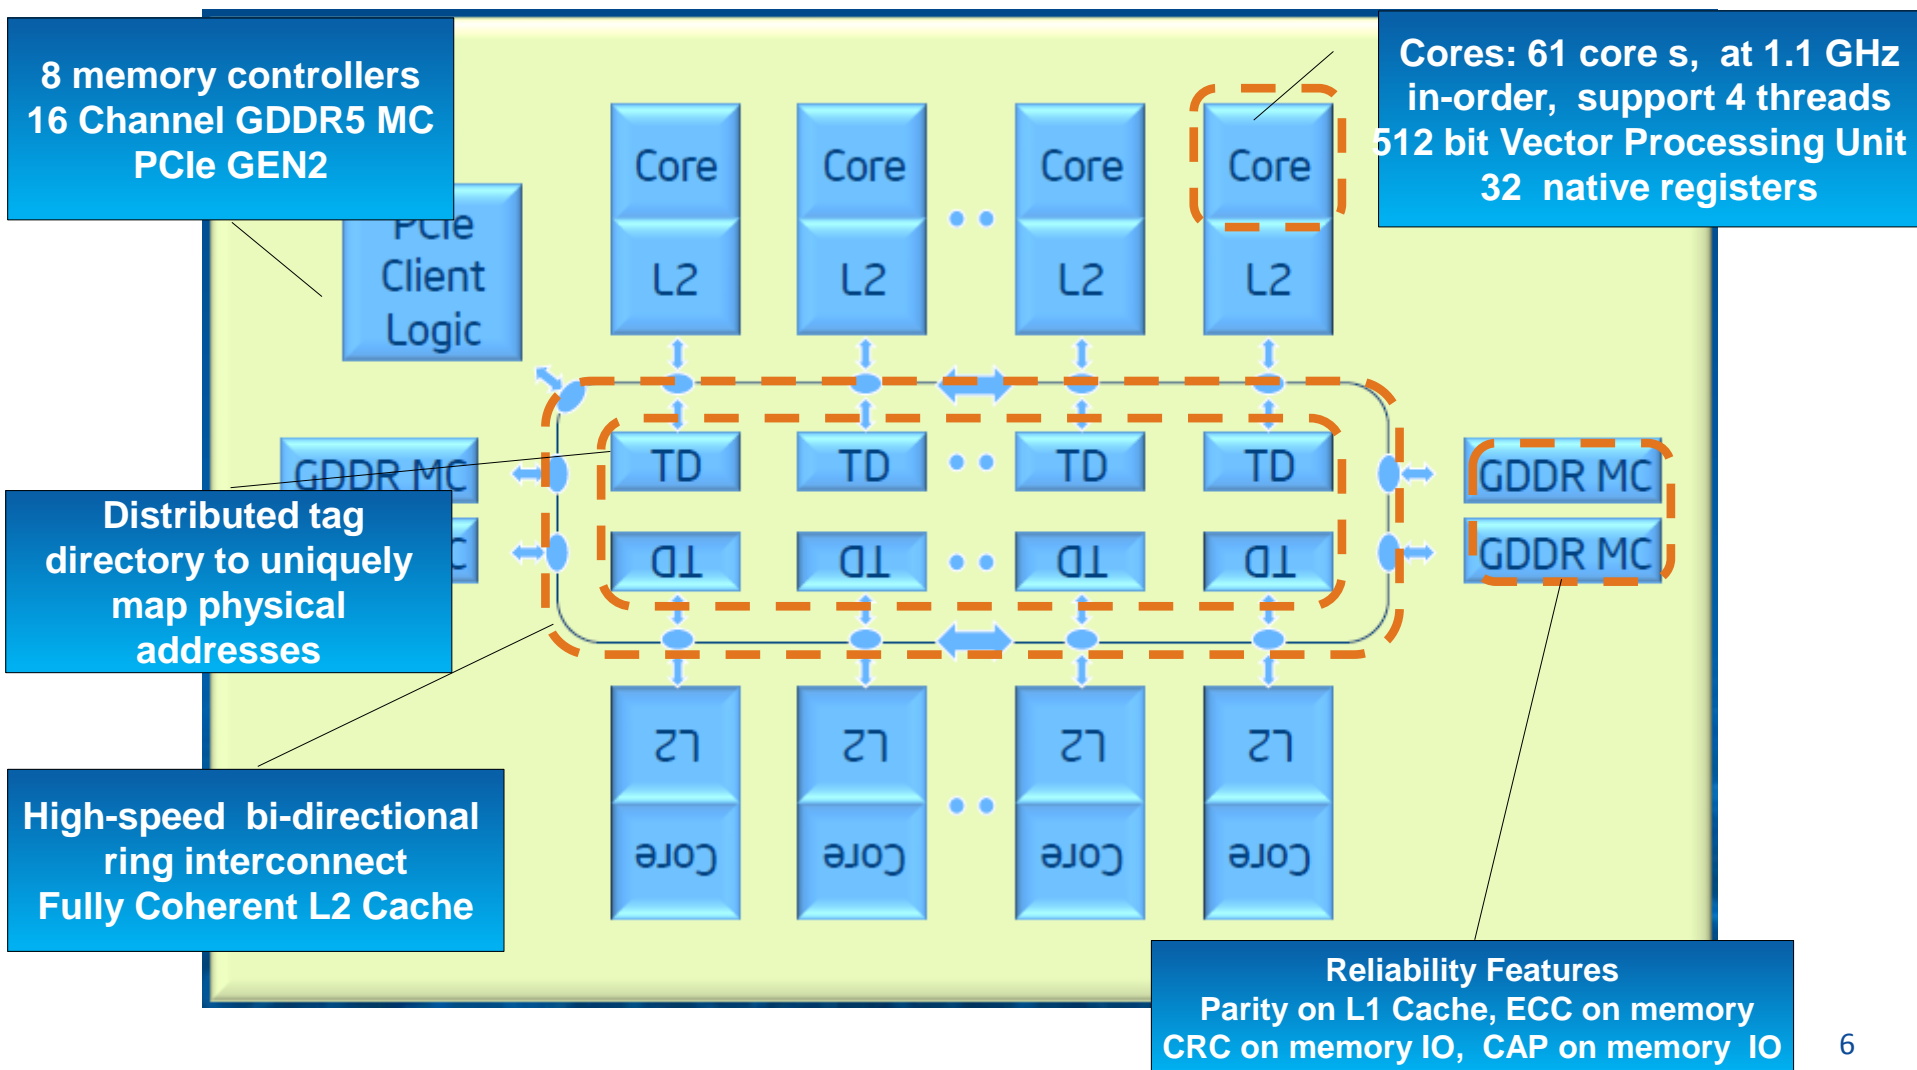


- Socket: mechanical component that provides mechanical and electrical connections between a microprocessor and a printed circuit board (PCB).
- QPI (Intel QuickPath Interconnect): high speed, packetized, point-to-point interconnection, that stitch together processors in distributed shared memory and integrated I/O platform architecture.



# Intel® Xeon Phi™ Architecture Overview

- Knights Core (KNC)



# Module Outline

---

- Intel Xeon and Intel® Xeon Phi™
- **System Software and OS**
- Intel KNL
- Programming Model

- Large SMP UMA machine – a set of x86 cores
  - 4 threads
    - ❑ 32 KB L1 I/D
    - ❑ 512 KB L2 per core
  - Supports loadable kernel modules
  - VM subsystem, File I/O
- Virtual Ethernet driver
  - supports NFS mounts from Intel® Xeon Phi™ Coprocessor
  - Support bridged network



# Intel® MIC Programming Considerations

- Getting full performance from the Intel® MIC architecture requires both a high degree of parallelism and vectorization
  - Not all code can be written this way
  - Not all programs make sense on this architecture
- KNC comes with 8GB or 16GB of memory
  - Only ~7GB or ~15GB is available to your program
    - ❑ The other ~1GB is used for data transfer and is accessible to your Intel® MIC Architecture code as buffers.
- Very short (low-latency) tasks not optimal for offload to the coprocessor
  - Costs that you need to amortize to make it worthwhile:
    - ❑ Cost of code and data transfer
    - ❑ Cost of process/thread creation
  - Fastest data transfers currently require careful data alignment

# Module Outline

---

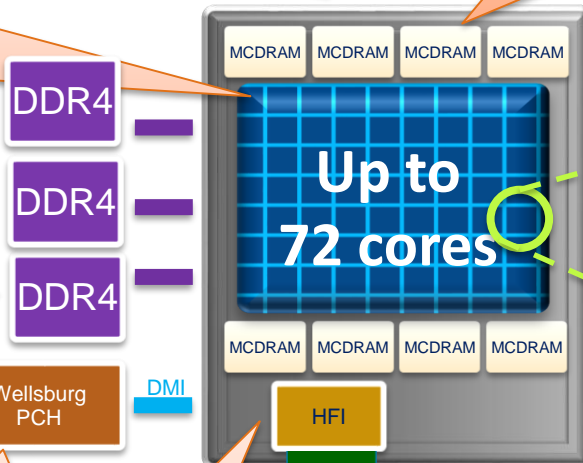
- Intel Xeon and Intel® Xeon Phi™
- System Software and OS
- **Intel KNL**
- Programming Model

# Knights Landing (KNL)


Over 3 TF DP peak  
Full Xeon ISA compatibility through AVX-512  
~3x single-thread vs. compared to Knights Corner

Up to 16GB high-bandwidth on-package memory (MCDRAM)  
Exposed as NUMA node  
~500 GB/s sustained BW

Up to 72 cores  
2D mesh architecture



2x 512b VPU per core  
(Vector Processing Units)



6 channels  
DDR4  
Up to 384GB

Common with  
Grantley  
PCH

2 ports  
**Intel Omni-Path Fabric On-package**  
50 GB/s bi-directional

Based on Intel® Atom Silvermont processor with many HPC enhancements  
Deep out-of-order buffers  
Gather/scatter in hardware  
Improved branch prediction  
4 threads/core  
High cache bandwidth  
& more


# Cluster modes

## One single space address

### Hemisphere:

the tiles are divided into two parts called hemisphere


Node 0



### Quadrant:

tiles are divided into two parts called hemisphere or into four parts called quadrants

Node 0




# Cluster modes

Cache data are isolated in each sub numa domain


## SNC-2:

the tiles are  
divided into two  
Numa Nodes




## SNC-4:

the tiles are  
divided into two  
Numa Nodes



# Integrated On-Package Memory Usage Models

## Integrated On-Package Memory Usage Models



Split Options:  
25/75% or 50/50%

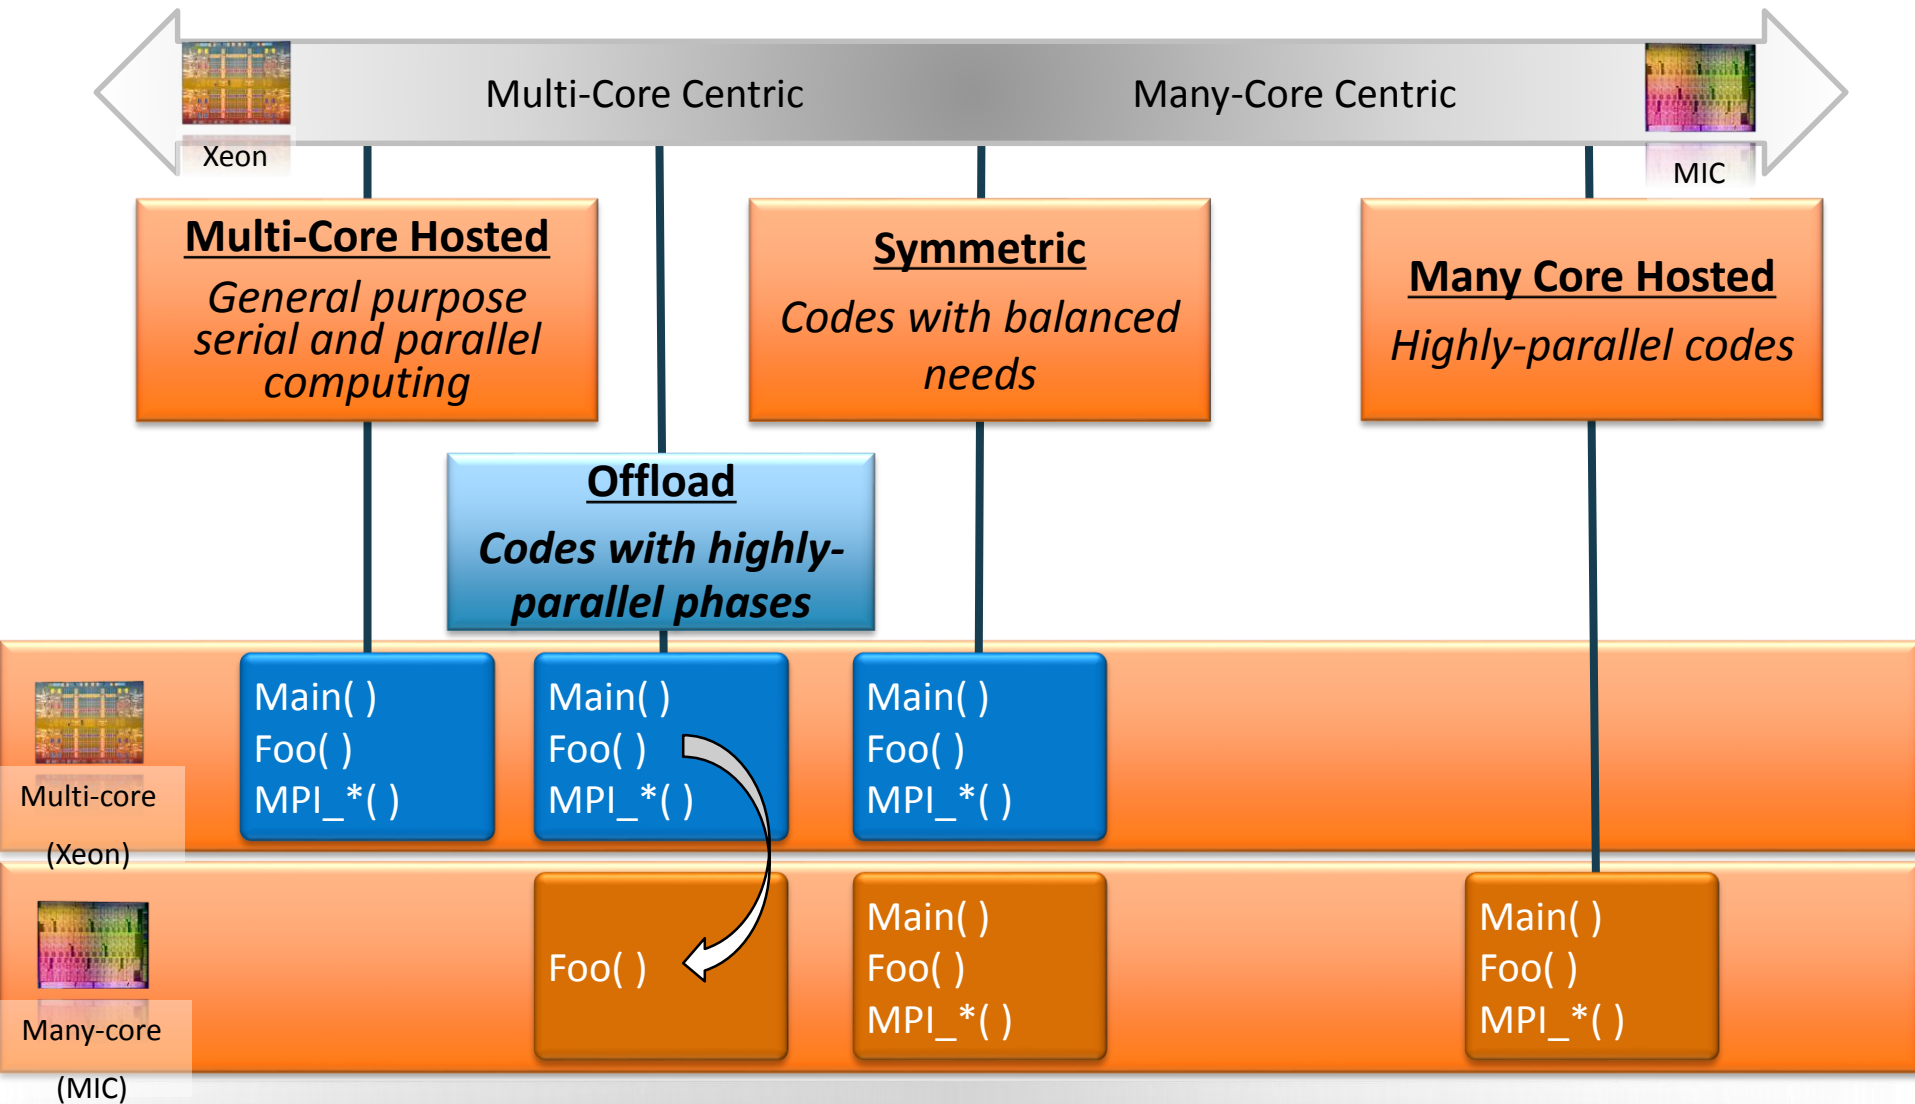
| Cache Model   | Flat Model  | Hybrid Model   |
|---|---|--|
| Hardware automatically manages the MCDRAM as a “L3 cache” between CPU and ext DDR memory  | Manually manage how the app uses the integrated on-package memory and external DDR for peak perf  | Harness the benefits of both Cache and Flat models by segmenting the integrated on-package memory  |
| <ul style="list-style-type: none"><li>▪ App and/or data set is very large and will not fit into MCDRAM</li><li>▪ Unknown or unstructured memory access behavior</li></ul> | <ul style="list-style-type: none"><li>▪ App or portion of an app or data set that can be, or is needed to be “locked” into MCDRAM so it doesn’t get flushed out</li></ul> | <ul style="list-style-type: none"><li>▪ Need to “lock” in a relatively small portion of an app or data set via the Flat model</li><li>▪ Remaining MCDRAM can then be configured as Cache</li></ul> |

# Module Outline

---

- Intel Xeon and Intel® Xeon Phi™
- System Software and OS
- Intel KNL
- **Programming Model**

# Programming Models





# Offload Programming Model Data Transfer

- Programmer designates variables that need to be copied between host and card in the offload directive using Pragma/directive;
- Variables and functions Allocation on both the host and device (C/C++):
  - `__attribute__((target(mic)))` [variable or function definition]
  - `__declspec(target(mic))` [variable or function definition] (windows only)
  - For entire files or large blocks of code
    - ❑ `#pragma offload_attribute (push, target(mic))`
    - ❑ `#pragma offload_attribute (pop)`

# Offload Programming Model Data Transfer

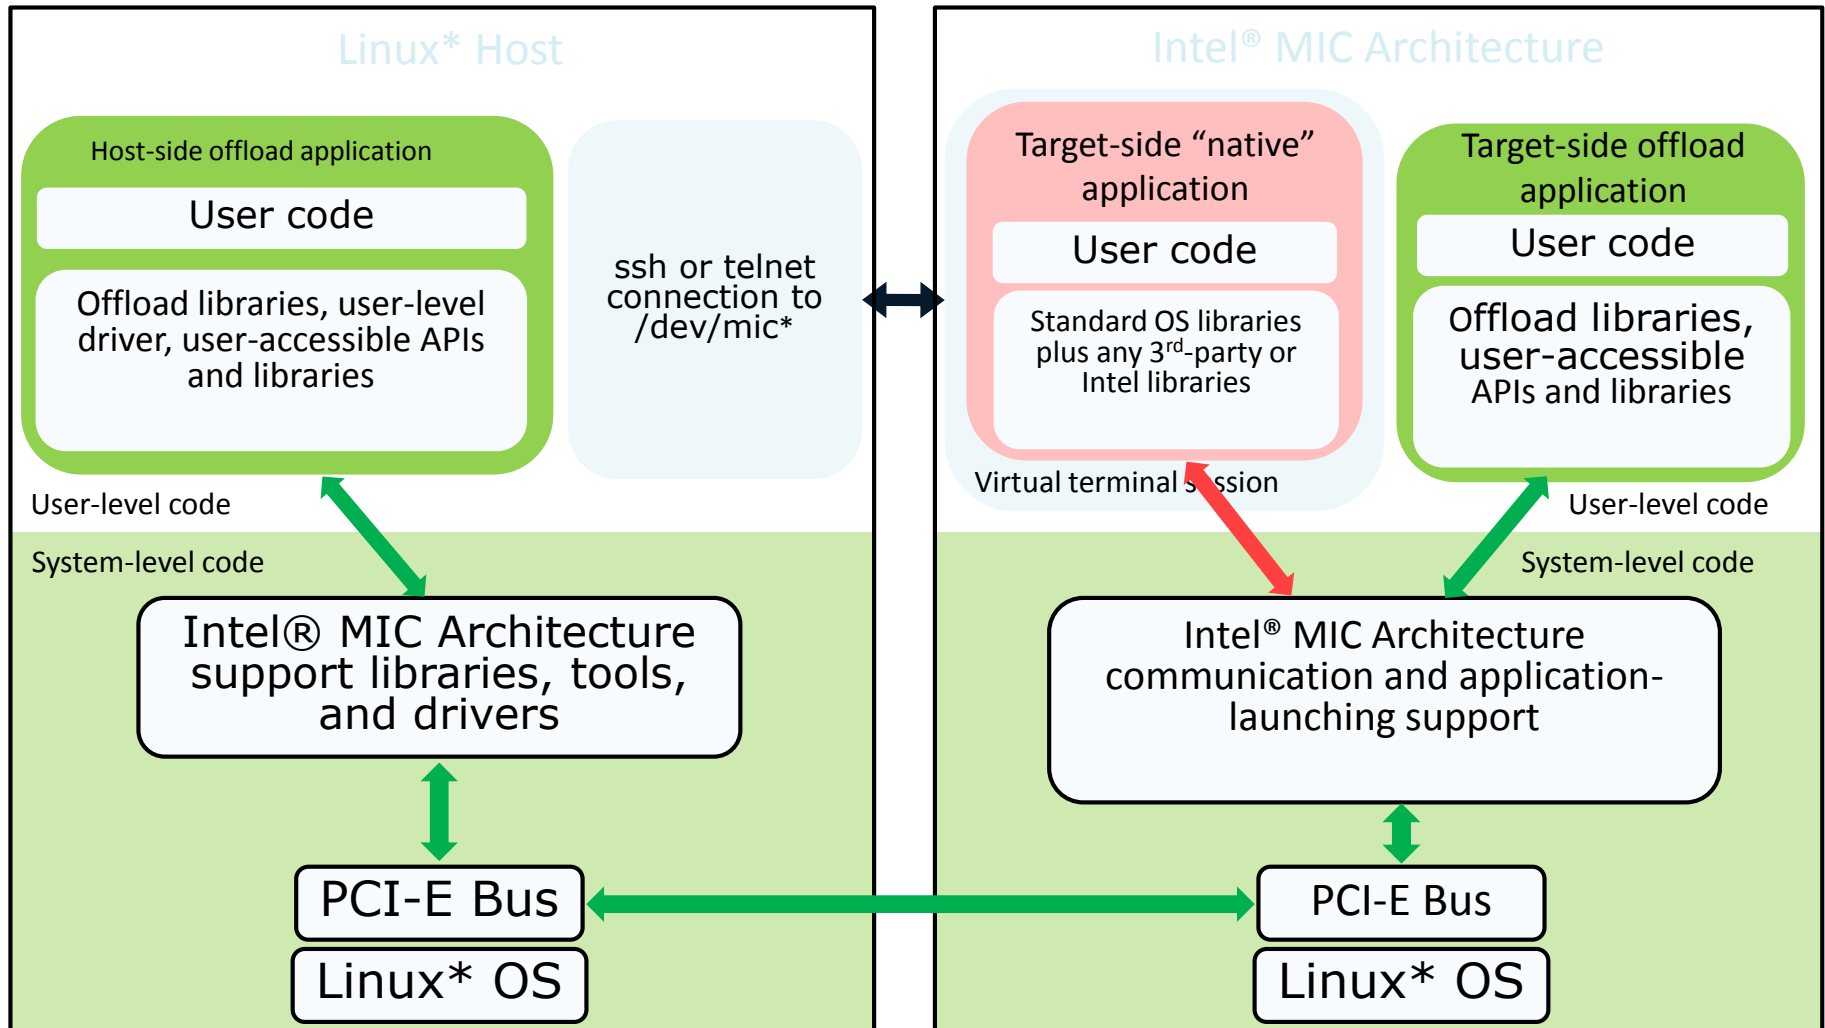
- Target
  - `#pragma offload target(mic[:dev-id]) [clauses]`
- dev-id: number of device to perform the offload
- clauses for explicit copy:
  - `nocopy` : allocate memory on device;
  - `in` : transfer a variable from host to device;
  - `out` : transfer a variable from device to host;
  - `inout` :
    - ❑ transfer a variable from host to device before start execution;
    - ❑ transfer a variable from device to host after finish execution;

# Offload Report

- **OFFLOAD REPORT:**
  - Measures the amount of time it takes to execute an offload region of code;
  - Measures the amount of data transferred during the execution of the offload region;
  - Turn on the report: `export OFFLOAD_REPORT=2`
- **[Var]** The name of a variable transferred and the direction(s) of transfer.
- **[CPU Time]** The total time measured for that offload directive on the host.
- **[MIC Time]** The total time measured for executing the offload on the target.
- **[CPU->MIC Data]** The number of bytes of data transferred from the host to the target.
- **[MIC->CPU Data]** The number of bytes of data transferred from the target to the host.

# Two Application Execution Environments

+Hybrid



# Examples

- Compiling
- Executing
- Offload
  - Code region
  - function
- Utilities
  - Micsmc
  - miccheck

